

E_{in}-style: Clothing Item Generation using GANs

Charlie Brayton (014559415)
Department of Software Engineering
San José State University
San José, California
charles.brayton@sjsu.edu

Mohit Patel (014501461)
Department of Software Engineering
San José State University
San José, California
mohit.patel@sjsu.edu

Andrew Selvia (014547273)
Department of Software Engineering
San José State University
San José, California
andrew.selvia@sjsu.edu

Dylan Zhang (013073437)
Department of Software Engineering
San José State University
San José, California
dylan.zhang@sjsu.edu

Abstract—The primary objective of this research is to classify and generate images of clothing items. It is a compilation of three distinct efforts: (1) to classify fashion-mnist [1] images and, inversely, add adversarial noise to test the classifier; (2) to generate novel images which emulate the true fashion-mnist images; (3) to infuse the grayscale images with life-like colors. Together, these endeavours pushed the team to explore topics at the forefront of modern machine learning, especially generative adversarial networks (GANs).

Index Terms—machine learning, computer vision, neural networks, generative adversarial networks

I. INTRODUCTION

The diverse branches of this project were made possible by the careful selection of a dataset. The fashion-mnist dataset was chosen specifically for its prevalence as a benchmark in recent research into neural networks. Its size, labels, and compatibility (with the long-established MNIST dataset) make it an ideal choice for evaluating deep neural networks (NNs) and GANs. Rather than attempt to introduce a truly groundbreaking NN architecture given the authors' limited experience, this project aims for breadth.

Each subproject explores a unique research area with industrial applications which may answer the associated questions:


- 1) Classification: Can we quickly identify an item of clothing to automate sorting and retrieval? Supposing such an automated system existed, how susceptible might it be to failure?
- 2) Generation: Can the product development lifecycle for clothing be accelerated by avoiding physical prototyping?
- 3) Colorization: Can grayscale clothing imagery be augmented with colors to enable A/B testing of styles to increase customer satisfaction?

To aid comprehension, each subproject is explored independently.

II. DATA

The fashion-mnist dataset is composed of 60,000 training images and 10,000 testing images. In keeping with the legacy

Fig. 1: Samples from fashion-mnist




of the traditional MNIST, fashion-mnist images are 28x28 grayscale pixels. Each image is labeled as one of the ten classes enumerated below:

- 0) T-shirt/top
- 1) Trouser
- 2) Pullover
- 3) Dress
- 4) Coat
- 5) Sandal
- 6) Shirt
- 7) Sneaker
- 8) Bag
- 9) Ankle Boot

TODO: reference Figure 1.

Fig. 2: Samples from fashion-mnist



III. IMPLEMENTATION

A. Classification

TODO: reference Figure 2.

Classification was performed using two neural network architectures, one using only dense layers and another using convolutional layers. The first architecture, shown in Figure 3, flattens the 28 by 28 pixel image into a vector 784 pixels long, which is then densely connected to a 128 node layer; finally, the 128 node layer is densely connected to a 10 node output layer.

The second architecture, shown in Figure 4, scans the image with a convolutional neural network using a 5 by 5 window and 64 layers; the results of the first convolutional layer are then scanned by another convolutional layer with a 5 by 5 window and 128 layers. The final step is to flatten the layer and densely connect it to the 10 node output layer.

The output layer of both architectures indicates the network's confidence that the image belongs to each of the classes, so a large value in 0th node and the 6th node would indicate that the network is uncertain whether an image should be classified as a T-shirt/top or just a Shirt. The neural networks described were constructed using Keras and TensorFlow in Jupyter notebooks.

The robustness of the classifier was tested by adding adversarial signals to the input images. The adversarial signal was calculated by determining the gradient of the loss function based on the input image used. This allows the adversarial

Fig. 3: First classifier architecture





Fig. 4: Second classifier architecture



signal to skew the image to the closest incorrect classification. Once the adversarial signal was determined it was scaled by an epsilon value (between 0 and 1) and added to the image; for adding this noise the image is considered to be all pixels that have a brightness value greater than 0. Limiting the noise to non-zero values means the classification isn't being skewed due to noise added in the black background portions of the image. Robustness was determined by comparing the accuracy of our classifiers against ever-increasing epsilon values.

B. Generation

A key goal from the onset of this project was to use GANs to generate images which are believable reproductions of those in fashion-mnist. During the planning stage, various projects which achieved the goal were identified as references. Two stood out above the rest: tensorflow-generative-model-collections [2] and pytorch-generative-model-collections [3]. Both take the same approach of training ten different GANs on various datasets, including fashion-mnist. The PyTorch implementation proved easier to execute on the SJSU HPC due to its more modern dependency stack, therefore it was chosen as the foundation of the generative component of this project.

The ten GANs explored for this project were:

- 1) Generative Adversarial Network (GAN; June 2014) [4]: The original GAN paper by Goodfellow, et al. introduced the world to the idea of training competing deep networks to improve generative models.
- 2) Conditional GAN (CGAN; November 2014) [5]: Followed an idea posed at the end of the original GAN paper to see how GANs might be used to fulfill conditional constraints such that the generated samples targeted a specific output class.
- 3) Information Maximizing GAN (infoGAN; June 2016) [6]: Maximizes mutual information during training to learn features without supervision which can be tuned to alter characteristics of the generated models (i.e. shifting the angle of MNIST-like digits).
- 4) Energy-based GAN (EBGAN; March 2017) [7]: An energy-based approach to GANs that yielded promising results for high-resolution images.
- 5) Auxiliary Classifier GAN (ACGAN; July 2017) [8]: Trains the generator to generate a classifier for each sample it produces, eventually proving GANs can learn over numerous diverse classes with high-resolution images. It achieves especially impressive results on ImageNet.
- 6) Least Squares GAN (LSGAN; April 2017) [9]: Swaps out the sigmoid cross entropy loss function for the least squares loss function to avoid vanishing gradients, thereby forcing the generator to produce higher quality images.
- 7) Wasserstein GAN (WGAN; January 2017) [10]: A theoretically-dense paper which introduces the critic into the standard game between a generator and discriminator to achieve more durable training and avoid mode collapse.
- 8) Boundary Equilibrium GAN (BEGAN; May 2017) [11]: Stabilizes training via an equilibrium, but more interestingly, explores the inverse relationship between diversity of generated images (modality) and image quality.
- 9) Wasserstein GAN - Gradient Penalty (WGAN_GP; December 2017) [12]: Enables training deeper, more complicated networks than previously possible by removing critic weight clipping from WGAN in favor of gradient penalties.
- 10) Deep Regret Analytic GAN (DRAGAN; December 2017) [13]: Calls into question the assumption that modality and quality are inversely correlated.

Each was trained independently on the fashion-mnist dataset for fifty epochs. The loss metric was tracked for each training session. To communicate the results of the training process more coherently, a snapshot of images produced by the generator is saved at the end of each epoch as demonstrated in Figure 8. Upon completion, each set of fifty images was compiled into a GIF to further streamline communication of the learning process.

C. Colorization

The images in the fashion mnist dataset are all grayscale images. We decided to build a machine learning model that

Fig. 5: Training samples for colorization




would fill colors into these black and white images intelligently. There were several techniques available to explicitly fill colors into the grayscale images. However, in order to smartly colorize the fashion mnist dataset we decided that Generative Adversarial Networks were the best choice for this task.

To build a model capable of smartly colorizing grayscale images we had to first build a dataset for training our GAN model. For this we combined the training and testing sets of the fashion mnist dataset. After this, we explicitly filled colors into these 70,000 grayscale images. Initially, the shape of each image was (28,28,1). After colorizing the images, their shape changed to (28,28,3). This marked the addition of 3 rows for each item in the dataset. This change happened to store the values of the red, blue and green elements of each image. Each label of the fashion mnist dataset was assigned different colors, for instance, ankle boots were colored brown and dresses were colored red. Now, we had 70,000 data instances each representing a colored fashion item from the fashion mnist dataset of the shape (28,28,3). Some instances of the explicitly colored training samples can be seen in Figure 5.

Once we had built our training samples, we decided to work on our GAN. We fed the generator in the network the original grayscale images from the fashion mnist dataset. The generator then tried to fill colors into these grayscale images of fashion items using some randomly set weights. These generated images were then passed on to the discriminator. Along with these images, the discriminator also received the explicitly colored images from the training sample. The discriminator tried to identify which image was created from the generator and which was drawn from the training sample. Based on the decision made on the classification task by the discriminator, the discriminator loss as well as the generator loss were computed.

The discriminator in the network learnt from the discriminator loss by backpropagating through the discriminator node and updating the weights associated with the discriminator. Similarly, the generator updated its weights using the results obtained from backpropagating through the generator node and using the generator loss. This process of updating the weights and learning was carried out iteratively. After running this iterative process for certain epochs, we had the final weights for the network that were capable of colorizing the items in

Fig. 6: DCGAN Architecture



the fashion mnist dataset intelligently.

We used DCGANs to define the neural network model. The DCGANs architecture displayed in Figure 6 gives a low level overview of the model used for colorizing the grayscale images of the fashion mnist dataset. It introduces the generators and discriminators along with the training process step by step.

DCGANs use a random noise vector as input, and the input is then scaled up into two-dimensional data similar to CNN's but in the opposite structure. For non-linear layers, we recommend using LeakyReLU for all layers for the discriminator. For the generator, we found that the BatchNormalize layer was the optimal choice, primarily because using BatchNormalize allows the CNN version of the generator to learn better, but it is impossible to use BatchNormalize for all layers. The output layer of the generator and input layer of the discriminator do not require a BatchNormalize layer.

Parameter setting for our project:


- batch size: 64
- Used Adam optimizer instead of Root Mean Square Propagation optimizer
- iteration times: 20000
- learning rate: 0.0001
- BatchNormalization: momentum=0.9

IV. RESULTS

A. Classification

Both classifiers had a training accuracy of 91%, and a test accuracy around 90%, 88% for the dense network and 90% for the convolutional network. This indicates we had some slight overfitting, but overall both classifiers performed well. The classifiers also seem to be fairly robust, since with an epsilon value of 0.05 we still saw a classification accuracy of 75%. Figure 7 shows the accuracy of our classifiers as epsilon is increased. The adversarial noise can be visually noticed with an epsilon of 0.1, so maintaining a 75% accuracy rate when the visual noise is largely unnoticeable to human observers is fairly robust. Based on the work of Goodfellow, Schlens, and Szegedy, the epsilon values at which adversarial noise effects a dataset is rather small [14]. Their classifier for the traditional MNIST data set produces an error rate of 89.4% with an epsilon of 0.25, and an error rate of 87% with an epsilon of 0.1. The difference in our accuracy compared to theirs is probably due to filtering the adversarial noise to only be on the image rather than in the black background as well. The main issues for the classifiers seem to be the similarity between some of the classes. The classifier has lower confidence in distinguishing between t-shirt and shirt, and between some

Fig. 7: Accuracy with adversarial noise



sandals and ankle boots. It could be interesting to see how the classifier changes if some of these groups are joined together.

B. Generation

The results of the original implementation were fully replicated. The presentation we delivered contains all of our results; a select few have been selected for this report.

After reading the ten papers in which these architectures are described, we could interpret the results more astutely. For instance, a common problem articulated by many of the papers is the challenge of training stability. They describe the fragility of the balance that exists between the discriminator and generator. Many of the papers specifically focus on ways to synchronize their learning, so the discriminator doesn't outpace the generator too early. The plots all tell that story; the discriminator's loss always improves faster than the generator. Notably, however, the research into training stability can be visually seen. In Figure 13, the more modern WGAN keeps the discriminator and generator in balance rather than allowing them to diverge as with the original GAN shown in Figure 11 or CGAN shown in Figure 12.

Another interesting observation manifest in the results is that GANs can learn to model specific classes. The original GAN does not constrain its generator to a particular output class, therefore it can be observed to shift its output based on the likelihood of fooling the discriminator in Figure 8; look specifically at the image directly to the right of the top-left corner (i.e. row 0, column 1) to see the generator switch from producing a Shirt to a Bag. In contrast, CGAN permits learning only on specified classes. Each column of Figure 9 holds predictions in a distinct output class.

Between Figure 8, Figure 9, and Figure 10, it seems the most recent (WGAN) actually produces the lowest quality predictions. Why? We suspect that its attempts to avoid mode collapse (the tendency for a GAN's predictions pool increasingly around a few high-confidence output classes) come at the cost of granularity. Of course, the DRAGAN authors would disagree. Ultimately, we need to perform additional research to understand this property of GANs better.

Fig. 8: GAN Generations




Fig. 9: CGAN Generations




Fig. 10: WGAN Generations




Fig. 11: GAN Loss Plot




Fig. 12: CGAN Loss Plot





Fig. 13: WGAN Loss Plot




C. Colorization

TODO: Cite [15]. TODO: Cite [16].

With the structure of the generative model and the discriminant model of the CNN structure provided in Figure 6 and training samples similar to Figure 5, the results achieved by our model can be seen in Figure 14.

It can be clearly seen that the generated images are not as good as the black and white examples, mainly due to the

Fig. 14: GAN colorized Clothing Images



difficulty in setting up the training set. We tried replacing too many background images and ended up using this new dataset. The generator must generate clothing images and have a valid background (retrieved from real images). With more training, additional training samples or some special techniques, one might be able to improve the generated images. However, since this is one way to use GAN to adapt to RGB images, we are delighted with the results.

For our project, we have used DCGAN to generate colorful clothing images. One can try to change the dataset to generate images that suits their personal need, or try to modify the neural network parameters to observe different generation effects. The size of non-trainable parameters in Adam optimizer was larger than RMSprop. This result excludes all parameters of the discriminator.

V. CONCLUSION

As evidenced by the diverse branches this project explored, neural networks are largely deserving of the hype they receive. We successfully leveraged them to perform multi-class classification and generative modeling over fashion-mnist. Furthermore, the breadth of GAN architectures provided broad exposure to modern research trends and problems GANs still face. No breakthrough theoretical discoveries were made. Nonetheless, our experiments and results lent us critical insight into the forces governing these techniques.

REFERENCES

- [1] H. Xiao, K. Rasul, and R. Vollgraf, "Fashion-mnist: a novel image dataset for benchmarking machine learning algorithms," 2017.
- [2] hwalsuklee, "tensorflow-generative-model-collections," 2017. <https://github.com/hwalsuklee/tensorflow-generative-model-collections>.
- [3] H. Kang, "pytorch-generative-model-collections," 2018. <https://github.com/znxlw/pytorch-generative-model-collections>.
- [4] I. J. Goodfellow, J. Pouget-Abadie, M. Mirza, B. Xu, D. Warde-Farley, S. Ozair, A. Courville, and Y. Bengio, "Generative adversarial networks," 2014.
- [5] M. Mirza and S. Osindero, "Conditional generative adversarial nets," 2014.
- [6] X. Chen, Y. Duan, R. Houthoofd, J. Schulman, I. Sutskever, and P. Abbeel, "Infogan: Interpretable representation learning by information maximizing generative adversarial nets," 2016.
- [7] J. Zhao, M. Mathieu, and Y. LeCun, "Energy-based generative adversarial network," 2017.
- [8] A. Odena, C. Olah, and J. Shlens, "Conditional image synthesis with auxiliary classifier gans," 2017.
- [9] X. Mao, Q. Li, H. Xie, R. Y. K. Lau, Z. Wang, and S. P. Smolley, "Least squares generative adversarial networks," 2017.
- [10] M. Arjovsky, S. Chintala, and L. Bottou, "Wasserstein gan," 2017.
- [11] D. Berthelot, T. Schumm, and L. Metz, "Began: Boundary equilibrium generative adversarial networks," 2017.
- [12] I. Gulrajani, F. Ahmed, M. Arjovsky, V. Dumoulin, and A. Courville, "Improved training of wasserstein gans," 2017.
- [13] N. Kodali, J. Abernethy, J. Hays, and Z. Kira, "On convergence and stability of gans," 2017.
- [14] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and harnessing adversarial examples," 2014.
- [15] C. Brayton, M. Patel, A. Selvia, and D. Zhang, "e-in-style," 2020. <https://github.com/AndrewSelviaSJSU/e-in-style>.
- [16] C. Brayton, M. Patel, A. Selvia, and D. Zhang, "Andrewselviasjsu/pytorch-generative-model-collections," 2020. <https://github.com/AndrewSelviaSJSU/pytorch-generative-model-collections>.